

MSc THESIS

LLVM-based ρ -VEX compiler

Maurice Daverveldt

Abstract

This thesis describes the development of a LLVM-based compiler for the ρ -VEX processor. The ρ -VEX processor is a runtime reconfigurable VLIW processor. Currently two compilers exist that target the ρ -VEX processor: a HP-VEX compiler and a GCC-based compiler.

We will show that both compilers have disadvantages that are impossible or very difficult to fix. That is why we have build a LLVM-based compiler that targets the ρ -VEX processor. The LLVM-based compiler can be parameterized in a way similar to the HP-VEX compiler. Further we will present certain optimizations that are new for LLVM-based compilers. These optimizations include a custom machine scheduler that avoids structural and data hazards in the generated binaries.

Finally we will demonstrate the operations of the LLVM-based compiler and compare the performance of generated binaries with the existing compilers. We will show that the LLVM-based compiler exceeds the performance and code quality of the GCC-based compiler. Binaries generated with the HP-VEX compiler outperform those of the LLVM-based compiler.

CE-MS-2014

LLVM-based ρ -VEX compiler asd

THESIS

submitted in partial fulfillment of the
requirements for the degree of

MASTER OF SCIENCE

in

COMPUTER ENGINEERING

by

Maurice Daverveldt
born in Utrecht, Netherlands

Computer Engineering
Department of Electrical Engineering
Faculty of Electrical Engineering, Mathematics and Computer Science
Delft University of Technology

LLVM-based ρ -VEX compiler

by Maurice Daverveldt

Abstract

This thesis describes the development of a LLVM-based compiler for the ρ -VEX processor. The ρ -VEX processor is a runtime reconfigurable VLIW processor. Currently two compilers exist that target the ρ -VEX processor: a HP-VEX compiler and a GCC-based compiler.

We will show that both compilers have disadvantages that are impossible or very difficult to fix. That's why we have build a LLVM-based compiler that targets the ρ -VEX processor. The LLVM-based compiler can be parameterized in a way similar to the HP-VEX compiler. Further we will present certain optimizations that are new for LLVM-based compilers. These optimizations include a custom machine scheduler that avoids structural and data hazards in the generated binaries.

Finally we will demonstrate the operations of the LLVM-based compiler and compare the performance of generated binaries with the existing compilers. We will show that the LLVM-based compiler exceeds the performance and code quality of the GCC-based compiler. Binaries generated with the HP-VEX compiler outperform those of the LLVM-based compiler.

Laboratory : Computer Engineering
Codenummer : CE-MS-2014

Committee Members :

Advisor: dr.ir. Stephan Wong, CE, TU Delft

Chairperson: dr.ir. K.L.M. Bertels, CE, TU Delft

Member: dr.ir. A. van Genderen, CE, TU Delft

Member: dr.ir. Guido Wachsmuth, CE, TU Delft

Dedicated to my family and friends

Contents

List of Figures	vii
List of Tables	ix
List of Acronyms	xi
Acknowledgements	xiii
1 Introduction	1
1.1 Motivation	1
1.2 Problem statement	3
1.2.1 Goals	4
1.3 Methodology	4
1.4 Thesis overview	5
2 Background	7
2.1 VEX System	7
2.1.1 Architecture	7
2.1.2 ISA	8
2.1.3 Run-time architecture	9
2.2 LLVM Compiler infrastructure	10
2.2.1 Front-end	11
2.2.2 LLVM IR	11
2.2.3 Code generation	12
2.2.4 Scheduling	13
2.3 Verification	14
2.4 Conclusion	14
3 Implementation	15
3.1 Tablegen	15
3.1.1 Register definition	16
3.1.2 Pipeline definition	17
3.1.3 Other specifications	18
3.2 Code generation	18
3.2.1 Instruction transformation	18
3.2.2 Instruction Lowering	19
3.2.3 Instruction selection	20
3.2.4 New instructions	20
3.2.5 Floating-point operations	22
3.2.6 Scheduling	22

3.2.7	Register allocation	22
3.2.8	Hazard recognizer	23
3.2.9	Prologue and epilogue insertion	23
3.2.10	VLIW Packetizer	23
3.3	New LLVM features	24
3.3.1	Generic binary support	24
3.3.2	Compiler parameterization	24
3.4	Conclusion	25
4	Optimization	27
4.1	Machine scheduler	27
4.2	Branch analysis	29
4.3	Generic binary optimization	30
4.3.1	Problem statement	30
4.3.2	Implementation	31
4.4	Large immediate values	31
4.4.1	Problem statement	32
4.4.2	implementation	32
4.5	Conclusion	32
5	Verification and Results	35
5.1	Verification	35
5.2	Benchmark results	36
5.2.1	General performance	37
5.3	Generic binary	39
5.4	Conclusion	42
6	Conclusion	45
6.1	Summary	45
6.2	Main contributions	46
6.3	Future work	46
	Bibliography	49
	List of Definitions	51
A	LLVM Quickstart guide	53
A.1	LLVM toolchain	53
A.2	XSTsim	53
B	LLVM Development guide	55
B.1	Building LLVM from source	55

List of Figures

1.1	MIPS pipeline [1]	1
2.1	ρ -VEX architecture	8
2.2	ρ -VEX instruction format	9
2.3	Basic compiler structure	10
2.4	Basic codegeneration process	13
3.1	<code>tblgen</code> instructions	16
5.1	ρ -VEX testbench	36
5.2	Absolute performance	39
5.3	Relative performance for increasing issue width	39
5.4	HP-LLVM relative performance	40
5.5	GCC-LLVM relative performance	40
5.6	Generic-Regular performance	42

List of Tables

2.1	ρ -VEX Register usage [2]	10
5.1	LLVM-based compiler performance in ns	41
5.2	HP-based compiler performance in ns	41
5.3	GCC-based compiler performance in ns	41
5.4	GCC-based compiler performance in ns	42
5.5	Register usage	42

List of Acronyms

VLIW Very Long Instruction Word
ILP Instruction Level Parallelism
OoO Out-of-Order Execution
CPI Clocks Per Instruction
LLVM Low Level Virtual Machine
GCC GNU Compiler Collection
ISD Instruction Selection DAG
DAG Directed Acyclic Graph
DFA Deterministic Finite Automaton

Acknowledgements

In 2007 I have started studying Electronic Engineering and Design at the University of applied sciences Utrecht. After graduating in 2011 I decided to persue a degree in Computer Engineering at the Technical University Delft. This thesis reflects the new things I have learned since beginning my study at the Electrical Engineering Department.

First I would like to thank my advisor Stephan Wong for giving me the opportunity and advising me during the writing of this thesis.

Secondly I would like to thank Rol Seedorf, Anthony Brandon and Joost Hoozemans for their discussions and help during the realization of this project. Their advice has proven to be priceless.

I would like to thank my friends and family for all the support they have given me over the years.

Maurice Daverveldt
Delft, The Netherlands
March 12, 2014

Introduction

In this chapter we will describe the reason and motivation for building a new compiler for the ρ -VEX processor. We will discuss the history of the ρ -VEX processor and of VLIW processors in general. Further we are going to see how a LLVM-based compiler can be an improvement over the current solutions that exist.

1.1 Motivation

In 2008 Thijs van As designed the first version of the ρ -VEX processor [3]. This processor uses a VLIW design and is based on the VEX ISA. The VEX ISA is a derivative of the Lx family of embedded VLIW processors [4] from HP/STMicroelectronics. Around this processor a set of tools has been developed in collaboration with the TU Delft, IBM, STMicroelectronics and other universities. Currently the ρ -VEX 2.0 tool suite include a synthesizable core, a compiler system and a processor simulator. A GCC based VLIW compiler has been developed by IBM.

A Very Long Instruction Word (VLIW) processor can execute multiple operations during a single clock cycle. A compiler is required to find parallelism between instructions and to provide scheduling that enables the VLIW processor to execute multiple operations during a single cycle.

A regular RISC type processor, such as the MIPS and ARM processor, contain a single instruction pipeline that executes instructions. Figure 1.1 shows a basic MIPS integer pipeline. By introducing pipelining registers the clock frequency of a processor can be increased because execution of an instruction is broken up into smaller and simpler parts. A pipeline can contain multiple instructions that are in different stages of execution.

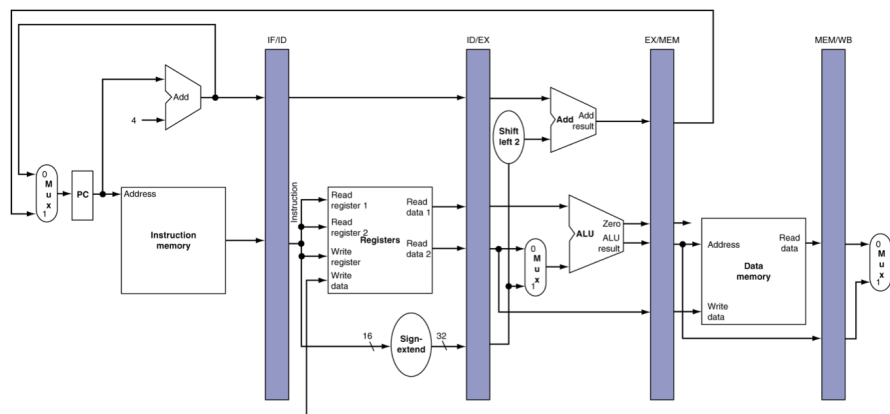


Figure 1.1: MIPS pipeline [1]

Introducing pipelining will generally lower the Clocks Per Instruction (CPI) rate of a processor to below 1.0. Special hardware has been developed, such as forwarding units, branch predictors, and speculative execution that will try to increase the CPI to a value that approaches 1.0.

If a higher than 1.0 CPI is desired multiple instructions need to be executed during a single clock cycle. Machines that can execute multiple instructions are called multi-issue machines. These types of processors use special hardware to find dependencies between instructions and to determine which instructions can be executed in parallel. These techniques include Tomasulo's algorithm for *Out of Order* execution and register renaming. Most modern processors use these techniques to increase performance.

Finding dependencies between instructions becomes increasingly complex when the issue width of machines is increased. The Pentium 4 processor demonstrated the limitations of further ILP extraction in a spectacular way. It used a 20-stage pipeline [5] with seven functional units. It operated on RISC-like micro-ops instead of x86 instructions and could handle 50 in-flight instructions at a time. The amount of silicon and energy that was dedicated to finding and executing ILP made the Pentium 4 processor very inefficient. The clock frequency increase that Intel expected the hyper pipelined processor (6-7 GHz) to deliver never materialized and the Pentium 4 Netburst architecture was dropped for a much simpler architecture.

In [6] the actual limitations of ILP extraction in hardware has been demonstrated and it has been shown that other techniques need to be used to find and execute more ILP.

VLIW processors differs from multiple issue machines in that parallelism is found during compile-time instead of during run-time. This results in a processor that can be made significantly simpler because the ILP extraction algorithms do not need to be implemented and because dependency checking is not required during run-time. Additional ILP can also be found with the compiler because the compiler has got a higher level view of the code that is to be executed. Optimizations such as swing modulo scheduling and loop vectorization are nearly impossible to achieve in hardware because the higher level structure of a program is no longer available. A compiler can interpret the higher level structure of a program and optimize the output for better scheduling.

The origins of the VEX ISA can be traced to the company Multiflow and John Fisher, one of the inventors of VLIW processors at Yale University [7]. Multiflow designed a computer that used VLIW processors to execute instructions up to 1024-bits in size. Along with these computers Multiflow also designed a compiler system that used trace based scheduling to extract ILP from programs. Reportedly the code base for the Multiflow compiler has been used in modern compiler such as Intel C Compiler (ICC) and HP VEX compiler because of the robustness and the amount of ILP that could be exposed by the compiler [8].

John Fisher has designed the VEX ISA as an example of VLIW type processors [9]. His work includes the design of an ISA, processor design, and a compiler system that generates code for this processor. Thijs van As developed a processor called ρ -VEX that is binary compatible with the VEX processor.

Currently two different compilers exist that target the ρ -VEX processor: the HP-VEX compiler and a GCC port developed by IBM. We will show that both existing

compilers are not optimal and that a new compiler is required for the ρ -VEX project. Further we will present a LLVM based compiler that targets the ρ -VEX processor with performance and features similar to the HP-VEX compiler.

1.2 Problem statement

Currently, both the HP-VEX and GCC compilers can be used to generate code for the ρ -VEX processor. Both compilers have got a number of advantages and disadvantages that will be explored. The compilers will be judged on the following subjects: Code quality, support, languages support, backend supported and customization possibilities.

HP-VEX:

- **Code quality:** Excellent code quality and ILP extraction.
- **Support:** Bad, no active community.
- **Front-end:** Bad, only support for C.
- **Back-end:** Not applicable since compiler is specifically targeted to one architecture.
- **Customization:** Customization possible through machine description. Further research on optimization strategies not possible because compiler is proprietary and closed source. Because of this expanding the functionality of the compiler is impossible.

GCC:

- **Code quality:** Excellent code quality with performance approaching that of commercial compilers.
- **Support:** There is a very active development community around GCC.
- **Front-end:** GCC supports a large number of programming languages including C, C++, Fortran and Java
- **Back-end:** Support exists for a large number of processors including x86, ARM, MIPS and ofcourse VEX
- **Customization:** Because GCC is open source the compiler can be customized to support new passes, optimizations and instructions.

Unfortunately GCC has a number of disadvantages that need mentioning.

- **VEX code quality:** The VEX backend for GCC has not been optimized and the quality of the code is quite low. Performance of GCC executables is lower then code compiled by the HP-VEX compiler. Some programs do not function correctly when compiled by GCC. Some programs are unable to be compiled by GCC.

- **VEX reconfiguration:** The current GCC VEX compiler does not support run-time reconfiguration. The compiler has been set to a 4 issue width ρ -VEX and this cannot be changed without rebuilding GCC.
- **Bloated:** GCC consists of millions of lines of code and is arguable one of the most complex programs in existence. This makes understanding GCC and developing for GCC very hard.
- **Complexity:** GCC is written in C. Design is complex, not very modular and documentation is not very good. Different parts of the compiler are linked in a complex way and it is very difficult to obtain a general overview on how the compiler operates. Because of the complexity it is difficult to achieve high performance in GCC.

The comparison shows that both the HP-VEX and GCC compilers have serious disadvantages. The fact that HP-VEX cannot be customized excludes it from further development for the ρ -VEX project. Bringing the GCC compiler performance and features up to the same level as HP-VEX will be very difficult because of the complexity involved with GCC development.

In 2000 the LLVM project has been started with the goal of replacing the code generator in the GCC compiler. LLVM provides a modern, modular design and is written in C++. Originally the GCC front-end was used to translate programs into LLVM compatible intermediate representation. Around 2005 the Clang project was started which aimed to replace the GCC front-end with an independent front-end that supports C, C++ and ObjC. Currently the LLVM based compiler offers performance that approaches GCC but offering a significant improvement in terms of modularity, ease of development and "hackability". In addition the LLVM compiler can also be used to target different architectures such as GPU's and VLIW based processors.

1.2.1 Goals

The main goal of this thesis is to develop a new compiler for the ρ -VEX system. The compiler will be based on the LLVM compiler. The new compiler should have the following characteristics:

- **Open source:** The compiler should be open source so the compiler can be customized and used for future research.
- **Code quality:** A new compiler should provide a significant improvement in terms of performance, code size and resource utilization.
- **Reconfigurability:** Characteristics of the ρ -VEX processor should be reconfigurable during run-time.

1.3 Methodology

The following steps need to be completed for successful implementation of a ρ -VEX LLVM compiler.

- Research ρ -VEX and VEX platform
- Research LLVM compiler framework
- Build LLVM based VEX compiler with following features:
 - 4 issue width VLIW
 - Code generation
 - Assembly emitter
- Add support for reconfigurability
 - VEX machine description
 - Reconfigure LLVM during runtime
- Optimize performance
 - Instruction selection
 - Hazard recognizer
 - Register allocator

1.4 Thesis overview

The thesis is organized as follows. In Chapter 2 we will discuss the architecture of the ρ -VEX processor and the workings of the LLVM compiler suite. This chapter will demonstrate the supported instructions, run-time architecture, and show the general architecture of the ρ -VEX processor. The chapter will also show how the LLVM compiler operates and what steps are involved during compilation.

In Chapter 3 will discuss how the ρ -VEX compiler was implemented. We will show how code is transformed from the LLVM Intermediate Representation into a ρ -VEX specific assembly language. In addition we will also discuss new functionality that has been added to the LLVM compiler.

Chapter 4 will discuss how the performance of the LLVM compiler has been optimized. Compilation problems that have been found are demonstrated and we will show how these problems have been resolved to increase performance of the binaries.

Chapter 5 will explore the performance of the new compiler. Performance will be compared to existing compilers in terms of issue width.

A conclusion and recommendations for future research is presented in chapter 6.

Background

In this section we will explore the background of the VEX system and of the LLVM compiler. This section will show the basic design of the ρ -VEX processor and how the ρ -VEX processor operates. Further we will also demonstrate the design of the LLVM compiler framework and how code is transformed from a high level language such as C/C++ to a target specific assembly language.

2.1 VEX System

The ρ -VEX processor is based on the VEX ISA [3]. The processor uses a VLIW architecture and is designed to serve as both a application-specific processor and a general-purpose processor. During synthesis the core can be reconfigured to alter the issue-width, the amount of physical registers, the type of functional units, and other parameters.

The VEX ISA defines hardware operations as syllables. An instruction is defined as a set of multiple syllables. The VEX operations are similar to 32 bit RISC operations.


2.1.1 Architecture

The architecture of the ρ -VEX core is shown in Figure 2.1 [10]. The core uses a five stage pipelined design. There are multiple functional units with different functionality, such as ALUs, Multipliers, Load / Store units, and Branch units.

The register file consists of 64 32-bit general purpose registers. These registers can generally be targeted by any instruction. In addition to the general purpose registers there also exists 8 1-bit Branch Registers. These registers are used by the branch operations to determine if a branch should occur.

The pipeline uses five stages to execute an instruction. For example, consider a ρ -VEX processor with a 4-issue organization:

- **Fetch stage:** An instruction is selected from the instruction memory using the Program Counter (PC) or the Branch Target address. Note that 1 instruction contains four different operations.
- **Decode stage:** Each operation is decoded in parallel and the needed registers are read from the register file.
- **Execute 1 stage:** In this stage the functional units produce results for ALU operations. In addition this stage is also used to write to the data memory when Store operations are executed. The ρ -VEX processor uses 16*32-bit multipliers that require two stages to produce a result. In this stage the multiplication operations are started.

Figure 2.1: ρ -VEX architecture

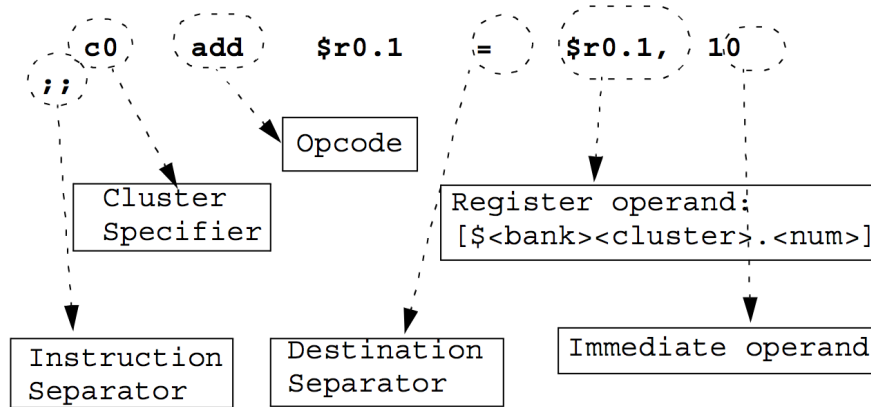
- **Execute 2 stage:** This stage produces results for MUL operations. In addition this stage reads data from the data memory when load operations are executed.
- **Writeback stage:** In this stage results that have been generated in the previous stages are committed to the register file.

The ρ -VEX processor uses a bypass network to forward operations to other pipeline stages when needed.

2.1.2 ISA

The assembly format for VEX instructions is shown in Figure 2.2 [2]. The destination of an operation is to the left of the "=" sign, while the source operands are listed on the right-hand side. On the right-hand side both registers and immediate values can be used as source operands. The ρ -VEX processor does not support multiple clusters and each instruction is executed on cluster 0.

The ρ -VEX processor supports 32-bit immediate values through a operations borrowing scheme. Each instruction can supports 8-bit immediate values but if larger values are required the adjacent operation is used to store the upper 24-bits of the immediate

Figure 2.2: ρ -VEX instruction format

value. This means that when using large immediate value the amount of operations that can be executed decreases.

Multiple classes of ρ -VEX instructions exists with the following properties:

- **Integer arithmetic operations:** These operations include the traditional RISC-style instructions such as ADD, SUB, AND, and OR.
- **Multiplication operations:** The VEX ISA defines multiple multiplication operations that use the builtin 16 * 32-bit multiplier. Operations include for example: Multiply Low 16 * Low 16, etc.
- **Logical and Select operations:** These operations are used to compare two registers to each other or to select between two values based on the result of a branch register. Operations include: CMPEQ, CMPNEQ, etc.
- **Memory operations:** Operations that load and store data from the data memory. Operations exist to store and load operands of different sizes such as LDW, LDH and LDB.
- **Control operations:** These operations are used to control the Program Counter of the ρ -VEX processor. Operations include: GOTO, CALL, BR, and RETURN.

2.1.3 Run-time architecture

The ρ -VEX Run-Time architecture defines the software conventions that are used during compilation, linking and execution of ρ -VEX executables. ρ -VEX programs are executed in a 32-bit environment where integers, longs and pointers are 32-bit values.

The following ρ -VEX register classes are used:

- **Scratch registers:** Caller-saved registers that are destroyed during function calls.
- **Preserved registers:** callee-saved registers that must not be destroyed during procedure calls.

Register	Class	Description
\$r0.0	Constant	Constant register 0
\$r0.1	Special	Stack-pointer: Holds the limit of the current stackframe. The SP is preserved across function calls.
\$r0.2	Scratch	Struct return pointer: If a function returns a struct or union the register contains the memory adres of the value being returned.
\$r0.3-\$r0.10	Scratch	Arguments and return values: Arguments that do not fit in the registers are passed using the main memory.
\$r0.11-\$r0.56	Scratch	Caller-saved scratch registers.
\$r0.57-\$r0.63	Preserved	Callee-saved registers that need to be preserved across function calls.
\$l0.0	Special	Link register: Used to store the return adres when a function call is performed.
\$pc0.0	Special	Program Counter
\$b0.0-\$b0.7	Scratch	Branch registers: Caller-saved registers.

Table 2.1: ρ -VEX Register usage [2]

- **Constant registers:** Contains a value that cannot be changed.
- **Special registers:** Used during call / return operations.

The 2.1 described the properties of all the available ρ -VEX registers.

2.2 LLVM Compiler infrastructure

LLVM is based on the classic three-stage compiler architecture shown in figure 2.3. The compiler uses a number language specific front-ends, an optimizer and target specific backends. Each module consists of a number of generic passes that are used to transform the code. This modular design enables compiler designers to introduce new passes and parts of the compiler without having to change the existing framework. Support for a new processor can be added by building a new back-end. The existing front-end and optimizer can be reused for the new compiler.




Figure 2.3: Basic compiler structure

The front-end is used to transform the plain text source code of a program into an

intermediate representation that will be used during compilation process. This transformation is achieved by performing the following steps:

1. **Lexical analysis:** Break input into individual tokens.
2. **Syntax analysis:** Using a grammar the sequence of tokens is transformed into a parse tree which represents the structure of the program. Clang uses a handwritten recursive descent parser for this transformation.
3. **Semantic analysis:** Semantic information is added to the parse tree, type checking is performed and a symbol table is built.

The resulting abstract syntax tree (AST) is transformed into LLVM IR and passed to the optimizer and backend of the compiler. These parts of the compilation process are completely language agnostic and do not require any other information from the front-end.

The optimizer is used to analyze and optimize the program. Optimization such as dead code elimination and copy propagation are performed during this phase but also more advanced operations that extract ILP (loop vectorization) can be enabled.

The back-end optimizes and generates code for a specific architecture. The LLVM IR is transformed into processor specific assembly instructions. In addition the backend also allocates registers, and schedules the instructions.

2.2.1 Front-end

The modular design of LLVM enables the compiler to be used as a part of the existing GCC compiler. For example, the dragonegg GCC plugin is designed to replace the GCC code generator and optimizer with the LLVM backend. This would enable LLVM to be able to use the existing GCC based front-ends and supported languages.

Clang has been developed to allow LLVM to operate independently of GCC. Clang is a front-end supporting C, C++ and ObjC. The front-end is designed to be closely integrated with the Integrated Development Environment (IDE) allowing more expressive diagnostic messages. In addition Clang also aims to provide faster compilation and lower memory usage [11].

2.2.2 LLVM IR

The front-end transforms a source code into the LLVM internal representation (LLVM IR). The LLVM IR is used to represent a high level language cleanly in a target independent way and is used during all phases of compilation. Instructions are similar to RISC instructions and can use three operands. Control flow instructions and type specific load/store instructions are used and an infinite amount of registers are available in Single Static Assignment (SSA) form. The LLVM IR is available in three different forms: human readable text, binary form, and an in-memory form [12].

The LLVM IR is designed to expose high-level information for further optimization. Examples of high-level information include dataflow analysis using the SSA form,

control-flow graphs, language independent type information and explicit use of pointer arithmetic.

Primitives such as voids, floats and integers are natively supported in the LLVM IR. The bitwidth of the integers can be defined manually. Pointers, arrays, structures and functions are derived from these basic types. The operations that are supported in LLVM IR are contained in the Instruction Selection DAG (ISD) namespace.

Object oriented constructs such as classes and virtual methods are not natively supported but can be built using the existing type system. For example a C++ class can be represented by a struct and a list of methods.

The SSA based dataflow form allows the compiler to efficiently perform code optimizations such as dead code elimination and constant propagation.

Figure 2.1 shows an example program in C. The equivalent LLVM IR representation is shown in figure 2.2.

```
int main() {
    int sum = 1;

    while(sum < 10)
    {
        sum = sum + 1;
    }
    return sum;
}
```

Listing 2.1: C example program

```
define i32 @main() nounwind ssp uwtable {
    %1 = alloca i32, align 4
    %sum = alloca i32, align 4
    store i32 0, i32* %1
    store i32 1, i32* %sum, align 4
    br label %2

; <label>:2                                ; preds = %5, %0
    %3 = load i32* %sum, align 4
    %4 = icmp slt i32 %3, 10
    br i1 %4, label %5, label %8

; <label>:5                                ; preds = %2
    %6 = load i32* %sum, align 4
    %7 = add nsw i32 %6, 1
    store i32 %7, i32* %sum, align 4
    br label %2

; <label>:8                                ; preds = %2
    %9 = load i32* %sum, align 4
    ret i32 %9
}
```

Listing 2.2: LLVM Intermediate representation

2.2.3 Code generation

During code generation the optimized LLVM IR is translated into machine specific assembly instructions. The modular design of LLVM enables generic algorithms to be used

for this process.

A backend is described in a domain specific language (DSL) called **tablegen**. The **tablegen** files describe properties of a backend such as available instructions, registers, calling convention and pipeline structure. During compilation of LLVM the **tablegen** files are converted into a C++ description of the backend. **tablegen** has been specifically designed to describe the backend structure in a flexible and generic way. Common features can be more easily described using **tablegen**. For example the **Add** and **Sub** instruction are almost identical and using **tablegen** can be described in a more generic way. This results in less repetition and reduces the chance of error in the backend description.

Because of the generic description of the backend large amount of code can be reused by each backend. Algorithms such as register allocation and instruction selection operate on the generic **tablegen** descriptions and do not require target specific hooks to operate correctly. An additional advantage of this approach is that multiple algorithms are available to achieve certain functionality. For example, LLVM offers the developer a choice between four different register allocation algorithms. Each algorithm has a number of advantages and disadvantages and the developer can choose between an algorithm which matches the target processor best.

At the moment not all parts of the backend can be described in **tablegen** and hand written C++ code is still needed. As LLVM matures more parts of the backend description should be integrated into the backend.

Figure 2.4 shows the basic codegeneration process. Each block can consist of multiple LLVM passes. For example the instruction selection phase consists of multiple passes that transform the input LLVM IR into a DAG that only contains instructions and types that are supported by the target processor.




Figure 2.4: Basic codegeneration process

2.2.4 Scheduling

The LLVM compiler uses basic blocks to schedule instructions. A basic block is a block of code which has exactly one entry point and one exit point. This means that no jump instruction exists with a destination in the block.

LLVM uses the **MachineBasicBlock** (MBB) class to represent a Basic Block. A MBB contains a list of **MachineInstr** instances. The **MachineInstr** class is an abstract way to represent instructions for the target processor.

Multiple MBB are used to create a **MachineFunction** instance. The **MachineFunction** class is used to represent a LLVM IR function. In addition to a

list of MBB the `MachineFunction` also contains references to the `MachineConstantPool`, `MachineFrameInfo`, `MachineFunctionInfo`, and `MachineRegisterInfo`. These classes keep track of target specific function information such as which constants are spilled to memory, the objects that are allocated on the stack, target specific function information, and which registers are used.

2.3 Verification

Verification of the LLVM compiler is extremely important. Performance of the generated binaries is irrelevant if only half of the binaries produce a correct result. The LLVM compiler will be verified by simulating the generated binaries.

Two simulators are available; `XSTsim` and `Modelsim`. `XSTsim` is an ISA simulator that can simulate a 4 issue width ρ -VEX processor. Output of the simulator is customizable and the simulator is fast enough to simulate large executables. `Modelsim` is used to perform a complete functional simulation of the ρ -VEX processor. The `Modelsim` simulation provides the highest accuracy because an actual ρ -VEX processor is simulated and is used for execution. The disadvantage of using `Modelsim` is the performance. Simulation of an executable will take a long time because the complete processor is simulated.

Verification of the compiler will be performed by writing test programs that generate certain instruction sequences. The output of these test programs will be compared to the expected result to check for errors in the backend.

Writing testprograms that have a high coverage of all the possible output patterns is impossible. Because of this a second round of verification will be performed by compiling benchmark programs that execute common programs. The output of the benchmarking programs will be compared to expected outputs to check for errors in the compiler.

2.4 Conclusion

This chapter discussed the background of the ρ -VEX processor and the basics of the LLVM compiler framework. Instructions that are supported by the ρ -VEX processor have been shown, the run-time architecture has been discussed and the register layout of the ρ -VEX processor has been discussed. The background information on the LLVM compiler should provide for a basic understanding on how the compiler operates and what parts need to be implemented to build a ρ -VEX backend.

Finally we have also discussed how the LLVM compiler will be verified. The verification step is extremely important to determine whether the binaries that are generated work correctly.

Implementation

The previous chapter demonstrated the architecture of the ρ -VEX processor and of the LLVM compiler. In this chapter we will show how the LLVM-based backend for the ρ -VEX processor has been implemented.

3.1 Tablegen

LLVM uses a domain specific language (DSL) called **tablegen** to describe features of the backend such as instructions, registers, and pipeline information.

tablegen uses a object-oriented approach to describe functionality. Information is described in classes and definitions that are called *records*. Inheritance is supported so classes can derive information from superclasses. In addition, multiclassses can be used to instantiate multiple abstract records at once.

The **tablegen** tool aims to provide a flexible way to describe processor features. For example processor instructions could be described as follows:

A class is created that represents an abstract instruction. The class will describe information that is of direct importance to code generation such as opcode, register usage and immediate values but also information that is needed during the code generation such as liveness information, instruction pattern and scheduling information.

```
class rvexInst<dag outs, dag ins, string asmstr, list<dag> pattern,
              InstrItinClass itin, Format f, CType type>: Instruction
{
}
```

The **rvexInst** class is used for all type of instructions that are supported by the ρ -VEX processor. Multiple instructions with common features such as **add**, **sub**, and **and** can be described in subclasses that inherit from the **rvexInst** class. For example, the class **ArithLogicR** holds arithmetic instructions that use three register operands. The class describes common features such as the instruction string format and the instruction pattern that is used during instruction selection.

```
class ArithLogicR<string instr_asm, SDNode OpNode,
                 InstrItinClass itin, RegisterClass RC, bit isComm = 0, CType ←
                 type>:
  rvexInst <(outs RC:$ra), (ins RC:$rb, RC:$rc),
            !strconcat(instr_asm, "\t$ra = $rb, $rc"),
            [(set RC:$ra, (OpNode RC:$rb, RC:$rc))], itin, type>
{
}
```

The following code shows how to define instructions that inherit from the **ArithLogicR** class. The instruction is defined as using the **ArithLogicR** class with

certain parameters that match the instruction properties. These properties include instruction string, LLVM IR opcode, and other information that is needed during compilation.

```
def ADD : ArithLogicR<"add ", add, IIALu, CPURegs, 1, TypeIIAlu>;
```

Figure 3.1 displays an example of how individual instructions inherit from higher level classes. The final class is a `rvexInstr` that contains a description of all the available ρ -VEX instructions.

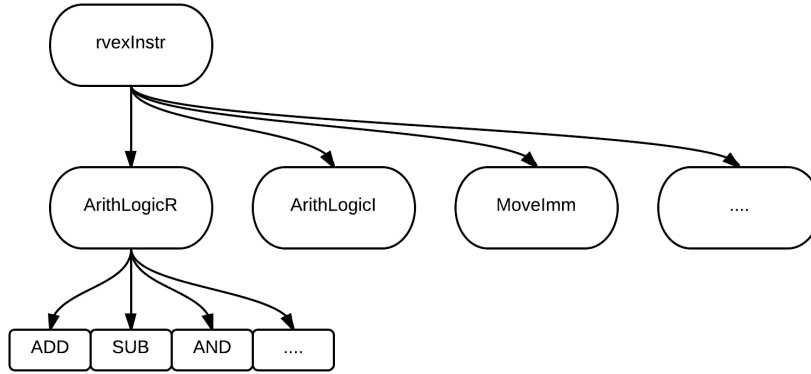


Figure 3.1: `tblgen` instructions

`tablegen` provides for a very flexible way to describe backend functionality. The existing LLVM backends use `tablegen` in a variety of ways which best match the target processor.

The `tblgen` tool is used to transform the `tablegen` input files into C++. The resulting C++ files contain enums, structs and arrays that describe the properties. The instruction selection part is transformed into imperative code that is used by the backend for pattern matching.

3.1.1 Register definition

LLVM uses a predefined class `register` to handle register classes. All ρ -VEX registers are derived from this empty class. The `rvexReg` class is used to define all type of ρ -VEX registers.

```
class rvexReg<string n> : Register<n> {
  field bits<7> Num;
  let Namespace = "rvex";
}
```

The `rvexReg` class is used to define the general purpose registers and the branch registers.

```
class rvexGPRReg<bits<7> num, string n> : rvexReg<n> {
  let Num = num;
}
```

```
class rvexBRReg<bits<7> num, string n> : rvexReg<n> {
  let Num = num;
}
```

Each physical register is defined as an instance of one of these classes. For example, R5 is defined as follows. The register is associated with a register number, a register string, and a dwarf register number that is used for debugging.

```
def R5 : rvexGPRReg< 5, "r0.5">, DwarfRegNum<[5]>;
```

The physical registers are divided in two register classes for the general purpose registers and for the branch registers. The register classes also define what type of value can be stored in the physical register.

```
def CPURegs : RegisterClass<"rvex", [i32], 32,
  (add
    (sequence "R%u", 0, 63),
    LR, PC
  )>;

def BRRegs : RegisterClass<"rvex", [i32], 32,
  (add
    (sequence "B%u", 0, 7)
  )>;
```

The branch registers have been defined to also use 32 bit values even though in reality the branch register is only 1 bit wide. This has been done because LLVM had a lot of trouble identifying the correct instruction patterns for compare instructions. The ρ -VEX compare instruction can produce results in both the CPURegs and the BRRegs as illustrated in the following example:

```
<1 bit>BRRegs = Operation, <32 bit>CPURegs, <32 bit>CPURegs
<32 bit>CPURegs = Operation, <32 bit>CPURegs, <32 bit>CPURegs
```

The LLVM compiler is unaware that when the compare instruction is used to define a 32-bit result only the lowest bit will be set. The compiler tried to resolve this issue by inserting `truncate` and `zero_extend` instructions even though this is not required. This has been solved by implementing the BRRegs as 32 bit wide so LLVM will not insert truncate and extend instructions when operating on these type of instructions.

3.1.2 Pipeline definition

`tablegen` can be used to describe the architecture of the processor in a generic way. LLVM will schedule an instruction to a processor functional unit during the scheduling pass. The following code describes the available functional units for a 4-issue ρ -VEX processor.

```
def P0 : FuncUnit;
def P1 : FuncUnit;
def P2 : FuncUnit;
def P3 : FuncUnit;
```

Each instruction is associated with an instruction itinerary. An instruction itinerary is used to group scheduling properties of instructions together. The ρ -VEX processor uses the following instruction itineraries.

```
def IIAlu          : InstrItinClass;
def IILoadStore    : InstrItinClass;
def IIBranch       : InstrItinClass;
def IIMul          : InstrItinClass;
```

The functional units and instruction itineraries are used to describe the properties of the ρ -VEX pipeline. The scheduling properties are derived from a description of an instruction stage with certain properties and the associated instruction itinerary. These properties include the *cycle count*, that describes the length of the instruction stage, and the functional units that can execute the instruction. The following itinerary describes a ρ -VEX pipeline with four functional units. Each functional unit is able to execute every instruction except for load / store instructions. Only P0 is able to execute load / store instructions. These instructions take two cycles to complete.

```
def rvexGenericItineraries : ProcessorItineraries<[P0, P1, P2, P3], [], [
  InstrItinData<IIAlu          , [InstrStage<1, [P0, P1, P2, P3]>]>,
  InstrItinData<IILoadStore    , [InstrStage<2, [P0]>]>,
  InstrItinData<IIBranch       , [InstrStage<1, [P0, P1, P2, P3]>]>,
  InstrItinData<IIMul          , [InstrStage<1, [P0, P1, P2, P3]>]>,
]>;
```

The machine model class is used to encapsulate the processor itineraries and certain high level properties such as issue width and latencies.

```
def rvexModel : SchedMachineModel {
  let IssueWidth = 2;
  let Itineraries = rvexGenericItineraries;
}
```

3.1.3 Other specifications

`tablegen` is also used to describe other properties of the target processor. LLVM has stated as goal to move more parts of the backend description to the `tablegen` format because `tablegen` offers such a flexible implementation. At the moment `tablegen` is also used to implement:

- Calling convention
- Subtarget features

3.2 Code generation

To understand how the compiler changes code from the LLVM IR representation to VEX assembly instruction it is necessary to understand how the code generation process works. The code generation process is divided into multiple steps, called passes, that are performed in order.

3.2.1 Instruction transformation

The instruction selection phase is completed in the following steps

- **Build initial DAG:** Transform the LLVM IR into a DAG that contains illegal types and instructions. The initial DAG is a one-to-one representation of the LLVM IR code.
- **Legalize instructions:** Illegal instructions are expanded and replaced with legal instructions.
- **Legalize types:** Transform the types used in the DAG to types that are supported by the target processor
- **Instruction selection:** The legalized DAG still contains only LLVM IR instructions. The DAG is transformed to a DAG containing target-specific processor instructions.

3.2.2 Instruction Lowering

The `rvexISelLowering` class gives a high level description of the operations that are supported by the target processor. The class can describe how the compiler should handle each LLVM IR instruction using four parameters: Legal, Expand, Promote or Custom. The default option is Legal which implies that the LLVM IR instruction is natively supported by the target processor.

3.2.2.1 Expanded instructions

The Expand flag is used to indicate that the compiler should try to expand the instruction into simpler instructions. For example consider the LLVM IR `UMUL_LOHI` instruction. This instruction multiplies two values of type `iN` and returns a result of type `i[2*N]`. Through expansion this instruction will be transformed into two `mult` instructions that calculate the low part and the high part separately.

```
setOperationAction(ISD::UMUL_LOHI, MVT::i32, Expand);
```

3.2.2.2 Promote instruction

Some instruction types are not natively supported and the type should be promoted to a larger type that is supported by the target processor. This feature is useful for supporting logical operations on Boolean functions. The following operation transforms an AND instruction that operates on a boolean value to a larger type.

```
setOperationAction(ISD::AND, MVT::i1, Promote);
```

3.2.2.3 Custom expansion

There are some instructions that cannot be expanded automatically by the compiler. To support these instructions the instruction expansion can be defined manually. For example consider the `MULHS` instruction that multiplies two numbers and returns the high part.

```
setOperationAction(ISD::MULHS, MVT::i32, Custom);
```

When a MULHS instruction is parsed the compiler will execute a function that describes the sequence of operations to lower this instruction. This sequence of instructions is implemented in the `LowerMULHS` function of the `rvexISelLowering` class. The `LowerMULHS` function is used to manually traverse the DAG and insert a sequence of instructions to support the operation.

For each instruction that requires custom lowering a `LowerXX` function has been defined.

3.2.3 Instruction selection

After instruction lowering the DAG contains LLVM IR operations and types that are all supported by the target processor but the DAG still contains only LLVM IR operations and no target-specific operations. The `rvexISelDAGToDag` class is used to match LLVM IR instructions to instructions of the target processor. The bulk of this class is generated automatically from the `tablegen` description but instructions can also be matched manually.

3.2.4 New instructions

The ρ -VEX processor supports some instruction that have no equivalent LLVM ISD operation. These instructions include `divs`, `addcg`, `min`, `max`, and others. Two stages are required to add support for these operations:

1. **Extend ISD namespace:** The new instructions will be added to the ISD namespace. This means that the LLVM IR will be *extended* with the new instructions.
2. **Instruction lowering:** Describe when these instruction should be inserted in the LLVM IR. For example, lowering of the LLVM IR `div` instruction uses a custom lowering function to describe the algorithm that uses the ρ -VEX `divs` instruction.
3. **Instruction matching:** New pattern matching rules need to be defined that map a custom instruction from the extended ISD namespace to the final target-specific ρ -VEX instructions.

For this example we are going to consider the ρ -VEX `divs` instruction.

3.2.4.1 Extend ISD namespace

The ISD namespace can be extended with custom by defining the instruction type and the instruction name. Because the `divs` instruction uses five operands a custom instruction type will be used. The following code shows the definition for the custom instruction type. This type describes an instruction that produces two results and consumes three operands.

```
def SDT_rvexDivs          : SDTypeProfile<2, 3
                          [SDTCisSameAs<0, 2>,
                           SDTCisSameAs<0, 3>,
                           SDTCisInt<0>, SDTCisVT<0, i32>,
                           SDTCisSameAs<1,4>,
                           SDTCisInt<1>, SDTCisVT<1, i1>]>;
```

The next step involves defining the name of the custom instruction. The instruction is defined as a custom `SDNode` and uses the instruction type that has been defined earlier. This operation expands the LLVM ISD namespace with custom operations that are only available in the ρ -VEX backend.

```
def rvexDivs              : SDNode<"rvexISD::Divs", SDT_rvexAddc>;
```

3.2.4.2 Instruction lowering

The `rvexISelLowering` class is extended with functions that produce the new ISD operation. For this example LLVM IR `div` instruction is custom lowered in the `LowerDIVS` function. In this function an algorithm is implemented that uses the new `divs` `SDNode`.

After instruction lowering a DAG will have been produced that contains only legal ISD operations and the new ISD operations that have been defined earlier.

3.2.4.3 Instruction selection

The following code describes the instruction class that is used by the `divs` instruction. This class describes the instruction string, register class properties and certain scheduling properties of this instruction. Note that the instruction pattern is empty because the current version of `tablegen` has no support for instructions that produce multiple results. A custom pattern matching function will need to be implemented in the `rvexISelDAGToDAG` class.

```
class ArithLogicC<bits<8> op, string instr_asm, SDNode OpNode,
                  InstrItinClass itin, RegisterClass RC, RegisterClass BRRegs, ↵
                  bit isComm = 0, CType type>:
  FA<op, (outs RC:$ra, BRRegs:$co), (ins RC:$rb, RC:$rc, BRRegs:$ci),
    !strconcat(instr_asm, "\t$ra, $co = $rb, $rc, $ci"),
    [], itin, type> { // Note empty instruction matching pattern
  let shamt = 0;
  let isCommutable = isComm; // e.g. add rb rc = add rc rb
  let isReMaterializable = 1;
}
```

The last step is to define the properties of the custom instruction.

```
def rvexDIVS              : ArithLogicC<0x13, "divs ", rvexDivs, IIAlu, CPURegs, BRRegs, ↵
  1, TypeIIAlu>;
```

The `rvexISelDAGToDAG` class is used to define the custom instruction selection patterns. This class implements the pattern matching code that is generated from the `tablegen` description files. The class has a separate `select` function to match instructions that have no pattern matching rules defined.

The `select` function uses switch statements to select between custom `SDNodes`. The following function implements the pattern matching rule for the `divs` instruction that

replaces the extended ISD `divs` instruction with the ρ -VEX instruction `rvexDIVS`. The `rvexDIVS` instruction has been defined earlier in the `tablegen` description files.

```
case rvexISD::Divs: {
  SDValue LHS = Node->getOperand(0);
  SDValue RHS = Node->getOperand(1);
  SDValue Cin = Node->getOperand(2);
  return CurDAG->getMachineNode(rvex::rvexDIVS, dl, MVT::i32, MVT::i32,
                                LHS, RHS, Cin);
  break;
}
```

3.2.4.4 Other cases

Some ρ -VEX specific instructions, such as the `SHXADD` instructions, are easier to support and do not need custom lowering. These instructions are also defined with empty pattern matching rules so the compiler will never insert them automatically. However the `SHXADD` instruction is easier to support because we can also match an instruction to a sequence of instructions. The following code describes a rule to replace the sequence `ADD (LHS<<1), RHS` with `SH1ADD LHS, RHS`.

```
def : Pat<(add (shl CPURegs:$lhs, (i32 1)), CPURegs:$rhs),
      (SH1ADD CPURegs:$lhs, CPURegs:$rhs)>;
```

3.2.5 Floating-point operations

ρ -VEX processor does not natively support floating-point instructions. Instead software functions are used to execute FP operations. During instruction lowering FP operations are translated into library calls that will execute the instructions.

The LLVM compiler uses library functions that are compatible with the GCC Soft-FP library. The LLVM compiler-RT library is compatible with the GCC soft-FP library and is used for execution of FP instructions. Compiler-RT is a runtime library developed for LLVM that provides for these library functions.

3.2.6 Scheduling

During the scheduling pass the `SDNodes` are transformed into a sequential list form. Different schedulers are available for different processor types. For instance the register pressure scheduler will always try to keep the register pressure minimal which works better for x86 type processors. For the ρ -VEX processor a VLIW scheduler is used.

The list still does not contain valid assembly instructions. Virtual SSA based registers are still used and all the stack references do not reference true offsets.

3.2.7 Register allocation

During register allocation the virtual registers are mapped to available physical registers of the target processor. The register allocator considers the calling convention, reserved registers and special hardware registers during allocation. In addition the register allocator also inserts spill code when a register mapping is not available.

Liveness analysis is used to determine which virtual registers are used at a certain time. The liveness of virtual registers can be determined easily through the SSA based form of the input list. Multiple register allocation algorithms are available. All algorithms operate on the liveness information.

3.2.8 Hazard recognizer

The ρ -VEX processor has no way to recognize hazards or halt execution of code for a cycle. Because of this the correct scheduling of instructions is important for correct execution. Consider the following sequence of code:

```
ldw $r0.2 = 0[$r0.2] # Load from main memory
;;
add $r0.2 = $r0.2, 2 # Add 2 to register
;;
```

After execution the register r0.2 will contain an undefined value because the load instruction has not completed execution. The compiler needs to insert an instruction or nop between the load and add instruction for correct execution. The correct instruction sequence should be the following:

```
ldw $r0.2 = 0[$r0.2] ## Load from main memory
;;
                        ## Empty instruction
;;
add $r0.2 = $r0.2, 2 ## Add 2 to register
;;
```

The `ScheduleHazardRecognizer` is only able to resolve structural hazards, not data hazards. In chapter 4 we will describe how the hazard recognizer has been combined with a new scheduling algorithm to resolve both data and structural hazards.

3.2.9 Prologue and epilogue insertion

After the register allocation pass the prologue and epilogue functions can be inserted. The prologue and epilogue pass is used to calculate the correct stack offset for each variable. Code is inserted that reserves room on the stack and saves / loads variables from the stackframe.

3.2.10 VLIW Packetizer

The packetizer pass is an optional pass that is used for VLIW targets. The packetizer receives a list of sequential machine instructions that need to be bundled for VLIW processors.

The `tablegen` pipeline definition is used to build a DFA that represents the resource usage of the processor. The DFA can be used to determine to which functional unit an instruction can be mapped and if enough functional units are available.

The DFA representation is powerful enough to consider different properties of functional units. For example consider a 4-issue width VLIW processor but with only one unit supporting load / store operations. The DFA can model this pipeline and guarantee only one load / store instruction will be executed per clock cycle.

The VLIW packetizer also checks if certain instructions are legal to bundle together. The packetizer can be customized to check for hazards such as data dependency hazards, anti dependencies and output dependencies. Custom hazards can also be inserted to make sure that control flow instructions are always in a single bundle.

3.3 New LLVM features

This section describes features that have been added to the LLVM compiler. Currently the ρ -VEX backend is the only backend that supports these kind of features.

3.3.1 Generic binary support

By disallowing RAW hazards in instruction bundles, binaries can be generated that can execute on any ρ -VEX processor irrespective of issue width. For example consider the following instruction packet. The RAW hazard will occur on the r0.8 register.

```
add $r0.10 = $r0.10, $r0.11
add $r0.12 = $r0.12, $r0.13
add $r0.14 = $r0.14, $r0.8
add $r0.8  = $r0.8, $r0.9
;;
```

Execution of this code on a 4-issue width processor will be fine but if this instruction is executed on a 2-issue width processor a problem will arise. The contents of register r0.8 will be changed during execution of the first bundle and execution of the second bundle will produce an invalid result.

This bundle should be split into two instructions during compilation. This can be achieved by checking for RAW hazards inside a bundle during the VLIW packetizer pass. The VLIW packetizer will detect these hazards and split the instructions into different bundles so correct execution is guaranteed.

3.3.2 Compiler parameterization

The HP VEX compiler supports a machine description file that describes properties of the processor. Using this machine description certain parameters, such as issue width, multiply units, load / store units, and branch units can be customized.

Currently the LLVM compiler supports parameterization through subtarget support. Backend features can be enabled and disabled with command line parameters. The ARM backend uses this approach to select between different ARMv7 architectures, floating point support and div/mul support. This approach would not be useful for the ρ -VEX processor because of the amount of features that can be customized. Each customizable feature would need a new subtarget. For example, when four features can be customized (W, X, Y and Z), then $W * X * Y * Z$ subtargets would be needed. Clearly this approach is not usable for the ρ -VEX processor where multiple parameters can be customized with a wide range of possibilities.

A different approach is used that changes the target description during runtime of the compiler. The target processor features are described in the `rvexMCTargetDesc` class.

During runtime a machine description file will be read parsed and information from this machine description will be used to update the `rvexMCTargetDesc` class.

The following properties can be customized through the machine description file:

- **Generic binary:** disable RAW hazard check in VLIW packetizer
- **Width:** issue width of the processor
- **Stages:** describes all the available functional units an the amount of cycles an instruction fills a functional unit.
- **Instruction itinerary:** maps an instruction itinerary to a function unit. Also describes at what cycle the output contains a legal value.

The parameters are also used to generate the DFA that will track the resource usage. During the translation of the `tablegen` file an algorithm is used to generate a DFA from the available functional units. This algorithm has been implemented in the `rvexMCTargetDesc` class.

The location of the machine description file is passed to the `rvexMCTargetDesc` through command line parameters. Custom command line parameters can be implemented in LLVM using `cl::opt` templates.

Some parameters such as the toggling of generic binary support is not handled through the `rvexMCTargetDesc` file. These parameters are used to set global state flags that can be read at anypoint in the compilation process. The `Is_Generic_flag` is used during the `rvexVLIWPacketizer` pass and during the liveliness calculation pass.

The following code demonstrates the config file for a 4 issue width ρ -VEX processor with support for generic binaries. The `Stages` parameter uses three values. The first value describes the amount of clock cycles a instruction occupies a functional unit the second parameter is a bitmask for which functional unit can handle this instruction. The third paramater is currently not used.

The `InstrItinerary` parameter describes which type of instruction maps to which `Stage`. The first value describes which Stage an instruction is going to use, the second parameter describes when the result of the operation is available.

```
Generic      = 1;
Width       = 4;
Stages      = {1, 15, 1}, {1, 1, 1}, {2, 2, 1};
InstrItinerary = {1, 2}, {1, 2}, {3, 5}, {2, 3}, {1, 2}, {1, 2};
```

The compiler parameterization feature can be easily ported to other LLVM backends. The only part that has changed is the `rvexMCTargetDesc` class.

3.4 Conclusion

In this section we have shown how the ρ -VEX backend for the LLVM compiler has been implemented. A short description has been provided on how the `tablegen` features are used for description of the a backend. We have shown how the ρ -VEX processor has been described in `tablegen` and how the different passes have been implemented. All the phases that are involved with code generation have been discussed.

Support for floating point operations has been added by porting the LLVM floating point library to the ρ -VEX processor.

Certain features that are required for the ρ -VEX processor are not available in the LLVM compiler. Features such as a machine description file are new to LLVM and we have shown what changes have been made to the LLVM compiler to support machine description files. In addition we have shown how the backend has been updated to provide support for the ρ -VEX generic binary format.

We have also shown how the performance of the backend has been improved by implementing custom scheduling features and a custom register allocator.

Optimization

The previous chapter described how the basic ρ -VEX LLVM compiler has been implemented. In this chapter we are going to discuss certain optimizations to improve the performance of binaries that are generated with the LLVM-based compiler.

4.1 Machine scheduler

A Machine Instruction scheduler is used to resolve structural and data hazards. The hazard recognizer that has been described earlier only resolves structural hazards. This means that the hazard recognizer can keep track of the functional units but it cannot keep track of data hazards between packets. The machine scheduler operates before the register allocator and is used to determine register allocation costs of each virtual register. This information is used during the register allocation pass to select a better mapping of physical registers to virtual registers that avoids expensive spills for commonly used virtual registers.

The Hexagon Machine scheduler has been customized to provide support for the ρ -VEX processor. The original Hexagon Machine Scheduler was used to perform packetization and to provide register allocation hints to the register allocator. The ρ -VEX machine scheduler has been enhanced with support for resolving data and structural hazards.

The machine scheduler passes uses the VLIW packetization information to build temporary instruction packets. A register allocation cost metric is used to determine an optimal scheduling and packetization of instructions. The following ρ -VEX instructions can produce data hazards that need to be resolved with the machine scheduler:

- **Multiply instructions:** 2 cycles
- **Load instructions:** 2 cycles
- **LR producing instructions:** 2 cycles
- **Compare and branch operations:** 2 cycles

The machine scheduler operates in two phases: Building of an instruction queue and scheduling of the instruction queue. During initialization of the machine scheduler two queues are build: The pending instruction queue and the available instruction queue. The available queue contains all the instructions that are available to schedule before a structural hazard occurs. The scheduler will schedule instructions until the available queue is empty and checks for any remaining structural hazards. If this hazard exists a `nop` instruction will be inserted, otherwise nothing will be done. The machine scheduler

will then load new instruction from the pending queue to the available queue until a new structural hazard occurs.

The scheduling phase is used to keep track of data hazards that have been scheduled in instructions. The algorithm builds temporary packets and checks for data dependencies between each packet. If a data dependency has been found a `nop` instruction is inserted to resolve the hazard. Listing 4.1 displays the algorithm in pseudo that is used to determine data dependencies.

```
# Get instruction from queue
Candidate = GetCandidate(Available_Queue)

if (!Candidate)
    # No Candidate found, structural hazard
    ScheduleMI(NULL, isNooped = True)

# Check latencies for all predecessors of candidate
for Predecessor in Candidate.Predecessors
    # Get instruction latency and scheduled cycle
    Latency      = Predecessor.Latency
    SchedCycle   = Predecessor.SchedCycle

    if (Latency + SchedCycle > CurrentCycle)
        # Only 1 Noop per <def> of instruction
        if (Predecessor.Nooped == false)
            # Schedule Noop
            InsertNoop = True

        #
        Predecessor.Nooped = True
        Predecessor.SchedCycle--

# Add Candidate to current packet
Reserve(Candidate)

If (Packet.full)
    # If packet is full increase scheduling cycle
    CurrentCycle++;
    Packet.clear

# Schedule instruction with Noop if required
Schedule(Candidate, InsertNoop)
```

Some corner cases exist that have been manually implemented in the machine scheduler.

The ρ -VEX processor has a 1 cycle delay between defining a branch register and using a branch register. The scheduler has been customized to insert `nop` instruction proceeding each branch instruction. This solution is not optimal because the empty instruction could also be used to execute an instruction that is not dependent on the proceeding instructions. Unfortunately this proved impossible to implement in the current version of the LLVM compiler because it is not possible to specify a delay between specific instruction classes

The following instruction sequence illustrates the current solution:

```
c0  cmpgt    $b0.0, $r0.2, 9
;;
;;
c0  br       $b0.0, .BB0_3
```

```
;;
```

The following code uses select instruction instead of branch instruction and does not need the 1 cycle delay between instructions:

```
c0  cmpgt    $b0.0, $r0.2, 9
;;
c0  slct     $r0.1 = $b0.0, $r0.2, $r0.3
;;
```

After the machine scheduler pass the temporary instruction bundles are deleted but the instruction ordering remains. Between the machine scheduler and VLIW packetization the only pass that can insert new code is the register allocator when spill code is introduced. During the `ExpendPredSpill` pass the spill code is checked for data dependencies and additional `nops` are inserted when required.

The final packetization occurs during the VLIW packetizer pass. The VLIW packetizer detects `nop` operations and outputs these instructions as *barrier* packets. These packets do not contain any instructions.

4.2 Branch analysis

Analysis of assembly files generated with the LLVM-based compiler showed that branches were not handled correctly. Consider the following C code:

```
if (c)
    return 1;
else
    return 2;
```

This code will be roughly translated into the following assembly code:

```
br    $b0.0, .BB0_2
;;
goto  .BB0_1          ## Goto not required
;;
.BB0_1
add  $r0.3 = $r0.0, 1
;;
goto  .BB0_3
;;
.BB0_2
add  $r0.3 = $r0.0, 2
;;
.BB0_3
return
;;
```

The first `goto` operation is not required because it will jump to an adjacent block of instructions. A branch analysis pass has been developed that can recognize jumps to adjacent blocks and remove unnecessary `goto` instructions.

In addition LLVM hooks for `insertbranch`, `removebranch` and `ReverseBranchCondition` have been implemented that allow the `BranchFolding` pass to further optimize branches.

4.3 Generic binary optimization

Research has shown [13] that generic binaries incur a performance penalty because of inefficient register usage. Consider the previous generic binary example again:

```
add $r0.10 = $r0.10, $r0.11
add $r0.12 = $r0.12, $r0.13
add $r0.14 = $r0.14, $r0.8
add $r0.8  = $r0.8, $r0.9
;;
```

This instruction packet is not legal because of the RAW hazard that is caused by r0.8. The current approach to fix the RAW hazard is to split the instruction packet into two separate instructions.

```
add $r0.10 = $r0.10, $r0.11
add $r0.12 = $r0.12, $r0.13
add $r0.14 = $r0.14, $r0.8
;;
add $r0.8  = $r0.8, $r0.9
;;
```

Because more instruction packets are used the amount of ILP that is extracted decreases and the performance of the resulting binary decreases. Performance could be improved by using a different register for the last assignment of r0.8.

4.3.1 Problem statement

This kind of optimization poses a significant challenge for VLIW type compilers because the register allocation pass is executed before the VLIW packetization. This implies that the register allocator has no information about which operations will be grouped together and for which operations an extra register should be used.

A solution could be to perform VLIW packetization before register allocation is completed. This would allow the register allocation pass to determine if a RAW hazard occurs inside a packet and to assign an extra register if needed.

In practice this approach is not possible because between the register allocation pass and the VLIW packetizer pass other passes are run that can change the final code. For example the register allocator pass can insert spill code and the prologue / epilogue insertion pass inserts code related to the stack layout. Inserting new instructions into instructions packets that have already been formed is very ugly because *packet spilling* could occur where a packet that is already full needs to move instructions to the next packet, etc.

The new machine scheduler could be customized to retain bundling information for the register allocator. Due to the complexity of implementing this approach we have chosen to change the register allocator itself and to make register allocation less aggressive.

The liveness allocation pass determines when a virtual register is used. The register allocator uses this information to create a register mapping with minimal register pressure.

By increasing the live range of a virtual register it should be possible to force the register allocator to use more registers when multiple virtual registers are used consecutively. Consider again the previous example:

```
add $r0.10 = $r0.10, $r0.11
add $r0.12 = $r0.12, $r0.13
add $r0.14 = $r0.14, $r0.8
add $r0.8  = $r0.8, $r0.9
;;
```

This would be transformed into the following code where the final `r0.8` assignment is changed to an unused register.

```
add $r0.10 = $r0.10, $r0.11
add $r0.12 = $r0.12, $r0.13
add $r0.14 = $r0.14, $r0.8
add $r0.15 = $r0.8, $r0.9
;;
```

4.3.2 Implementation

The `LiveIntervals` pass is used to determine the live ranges of each virtual register. The pass uses the `LiveRangeCalc` class. Each virtual register has an associated `SlotIndex` which tracks when the register becomes live and when the register is killed. The `ExtendToUses` method of the `LiveRangeCalc` class is used to update the `SlotIndex` to match the latest use of a virtual register. The `SlotIndex` class also provides method that give information on the `MachineBasicBlock` (MBB) in which the virtual register is used.

Extending of the liverange has been enabled by getting the boundary of the MBB from the `SlotIndex` and by extending the `SlotIndex` to this boundary. This enables the virtual register to be live for the duration of the basic block and will make sure the register allocator will not assign a new virtual register to a previously used physical register.

This approach is not optimal because it will increase the register usage even if RAW hazards do not occur. If more virtual registers are used then physical registers are available the execution speed will drop because extra spill code needs to be inserted.

4.4 Large immediate values

The ρ -VEX processor has support for using 32-bit immediate values. 8-bit immediate values can be handled in a single ρ -VEX instruction. Values larger than 8-bit values borrow space from the adjacent ρ -VEX instruction. The following code examples show the maximum amount of instruction in a packet for a 4-issue width ρ -VEX processor.

```
add $r0.10 = $r0.10, 200
add $r0.11 = $r0.11, 200
add $r0.12 = $r0.12, 200
add $r0.13 = $r0.13, 200
;;
```

```

add $r0.10 = $r0.10, 2000
add $r0.11 = $r0.11, 200
add $r0.12 = $r0.12, 200
;;

```

Large immediate values are used throughout the ρ -VEX ISA. Not only arithmetic instruction can use large immediate values but also load and store instructions to represent the address offset.

```

ldw $r0.2 = 2000[$r0.2]
;;

```

4.4.1 Problem statement

Large immediate values can be supported by creating a new instruction itinerary with support for large immediate values. This instruction itinerary would be special because it requires two functional units during VLIW packetization. The algorithm that builds the resource usage DFA needs to be updated to reflect instruction itineraries with multiple functional unit usage.

Unfortunately this approach is impractical for a couple of reasons. Each instruction that supports immediate values needs to be implemented twice, once for small immediate values and once for large immediate values. This would increase the risk of errors in the Tablegen files because each instruction has multiple definitions that need to be updated when changes are made.

A second approach is to update the VLIW packetizer and to recognize large immediate values during the packetization pass.

4.4.2 implementation

During the packetization pass each instruction that uses an immediate value is checked before it is bundled in an instruction bundle. The packetizer determines the size of an immediate value and reserves an extra resource in the DFA when a large immediate value is used.

4.5 Conclusion

In this section we have discussed the optimizations that have been implemented to increase performance of ρ -VEX binaries that have been generated with the LLVM compiler.

The `rvexMachineScheduler` pass is used to handle structural and data hazards. Temporary instruction packets are generated and are filled with instructions that have been selected using cost based scheduling algorithms to reduce register pressure. The pass enables ρ -VEX binaries to execute correctly and to perform better than binaries that have not been scheduled using the `rvexMachineScheduler`.

The branch analysis optimization is used to erase unnecessary `goto` statements from the code. In addition hooks have been provided that allow the LLVM `BranchFolding` pass to further optimize branches that are used in ρ -VEX binaries.

The generic binary optimization allows binaries with generic binary support to perform on par with regular binaries. The performance of generic binaries will only degrade once the register pressure becomes too high and spill code needs to be inserted.

The immediate value optimization allows more efficient use of available instructions of the ρ -VEX processor.

Verification and Results

The previous chapters described how the LLVM-based ρ -VEX compiler has been implemented. In this chapter we are going to verify the correct operation of the compiler and we are going to measure the performance of the binaries that are generated with the LLVM-based compiler.

5.1 Verification

Unit testing has been used to verify the correct operation of the LLVM compiler. The tests have been performed using the XSTsim simulator. XSTsim can only print pipeline and register information to the terminal. It is not possible to parse strings or information from the program that is executing back to the user. To check if tests are executed correctly we check the return statement after execution of a benchmark. The value in the return register indicates whether the test executed correctly or where at which point the test failed.

- **arit.c:** Integer arithmetic tests for `char`, `short`, `int` and `long long`.
- **if.c:** Integer and boolean comparison operators.
- **float.c:** Testing of floating point library.
- **func.c:** Tests involving pointers and structures.
- **global.c:** Tests involving global integers, arrays and structures.
- **call.c:** Function calls.
- **func_pointer.c:** Function calls using function pointers.
- **loop.c:** Basic while loops.
- **misc.c:** Others tests.

During preliminary testing of the benchmark additional errors have been found. The verification tests have been updated to catch these errors. Unfortunately some benchmarking errors are not possible to define as a unit test. Some errors, such as scheduling and register allocation errors, only occur in complex programs. Translating these errors to simple unit tests is not possible because they depend on the higher-level structure of the program.

5.2 Benchmark results

The powerstone benchmark [14] has been used to evaluate the performance of the LLVM-based compiler. The benchmarks consist of a number of programs that test certain functionality. In addition to performance evaluation the benchmarks are also useful to check the executable correctness of the generated binaries.

Benchmarking can be performed using the architecture simulator XSTsim and using the hardware simulator. We have chosen to use the Modelsim hardware simulator because a complete logical simulation of the processor is used for evaluation. Modelsim builds a simulation environment using the ρ -VEX VHDL files. A testbench is used to generate test signals. For this simulation the testbench is used to load the instruction and data memory of the ρ -VEX processor and to generate correct clock and enable signals to start execution. In Figure 5.1 the testbench process is displayed.




Figure 5.1: ρ -VEX testbench

Modelsim wave viewer can be used to monitor execution of the processor. The wave viewer has been used to check the contents of the registers and to monitor the instruction pipeline when necessary.

The architectural simulator has not been used because it only operates on a description of the ρ -VEX processor. Bugs that are present in the architectural simulator are not necessarily present in the actual processor. By using the Modelsim simulator we can be sure to a reasonable extent that binaries that execute in Modelsim will also execute on actual hardware.

The following benchmarks have been used for evaluation:

- **adpcm:** codec for voice compression.
- **bcnt:** Bitwise shift and operations on 1K array.
- **blit:** Graphics application

- **compress:** UNIX compression utility
- **crc:** Cyclic redundancy check.
- **engine:** Engine control application
- **g3fax:** Group 3 fax decode.
- **matrix:** Matrix multiplication.
- **pocsag:** Communication protocol for paging applications.
- **ucbqsort:** Quicksort algorithm.

Unfortunately not all benchmarks execute correctly. Currently all benchmarks work on the XSTsim simulator but some benchmarks do not work correctly using the Mod-
elsim simulator. This indicates possible bugs in the LLVM-based compiler that are probably related to scheduling errors. XSTsim has proven to be less strict when executing instructions with different latencies. Problems have been found with the following benchmarks:

- **fir:** Unclear how program operates. Omitted because benchmark was also unable to compile on X86 based systems.
- **jpeg:** Output not correct. Output is very close to expected value so this probably indicates a minor error in the compiler.
- **qurt:** Infinite loop.
- **des:** No check on output so impossible to determine if execution was correct.
- **v42:** Output not correct.

5.2.1 General performance

General performance of all the compilers that target the ρ -VEX processor are shown in Tables 5.1, 5.2 and 5.4. As expected the HP-based compiler performs excellent but unfortunately is not able to produce correct executables for higher issue widths. The GCC compiler has only been implemented for a 4 issue width ρ -VEX processor.

The absolute performance of the compilers is displayed in Figure 5.2. As expected the HP-based compiler generates the best performing binaries. The LLVM-based compiler offers some speed improvements over the GGC-based compiler for most benchmarks.

The HP-based compiler generates excellent performing binaries because it has a superior scheduling techniques for finding and extracting ILP in source code. The HP-based compiler uses a trace based scheduling technique which enables better ILP extraction. In addition to this the HP-based compiler also seems to do more optimization. Even when compiling with `-O0` the compiler already performs certain optimizations that are not available for LLVM-based compiler. For example, compare the output for the following simple C program:

```

int main() {
    int a = 3, b = 2, c;

    c = a + b;

    return c;
}

```

Listing 5.1 displays the output of the HP-based compiler and Listing 5.2 shows the output of the LLVM-based compiler. The output shows immediately that the HP-based compiler has eliminated the add operation and has copied the final value of the operation straight to the return register. The LLVM compiler is only able to perform these kind of operations at higher optimization levels.

```

    add $r0.3    = $r0.0, 5      ## Move to return register
;;
    return
;;

```

Listing 5.1: HP compiler output

```

    add $r0.2    = $r0.0, 2
    add $r0.3    = $r0.0, 3
;;
    add $r0.3    = $r0.2, $r0.3  ## Move to return register
;;
    return
;;

```

Listing 5.2: LLVM compiler output

Close inspection of the absolute performance reveals some interesting facts related to the issue width of the processor. Increasing the issue width from 2 issue to 4 issue leads to a significant increase in performance. Further increasing the issue width to 8 issue machine does not lead to an increase in performance. This is shown in Figure 5.3. For some benchmarks such as `bcnt` the LLVM-based compiler is able to find extra parallelism but for most benchmarks the performance increase is non-existent.

Manual inspection of the assembly files that are generated shows that both compilers are able to generate instruction packets that use a higher number of functional units. These large instruction packets do not lead to an increase in performance because they are not contained inside loop structures. The amount of times large instruction packets are executed is limited.

Figure 5.4 shows the relative performance of LLVM-based binaries compared to HP-based binaries. Higher than 100% indicates LLVM-based binaries performing better than HP-based binaries. As expected HP-based binaries perform better than LLVM-based binaries. Inspection of the generated assembly files indicate that the HP-based compiler is able to fill significantly more functional units than the LLVM-based compiler. Further the LLVM-based compiler is overly aggressive in inserting nop instructions to reduce structural and data hazards.

Figure 5.5 shows the relative performance of LLVM-based binaries compared to GCC-based binaries. Performance of the LLVM-based compiler shows mixed results. Some




Figure 5.2: Absolute performance




Figure 5.3: Relative performance for increasing issue width

benchmarks perform significantly better such as `matrix` and `ucbqsort` but both `compress` and `pocsag` perform worse. Manual inspection of these benchmark do not show a big reason for decreased performance except for the overly aggressive `nop` insertion of the LLVM machine scheduler.

5.3 Generic binary

The performance of generic binaries has been tested by generating three sets of binaries for a 4 issue width ρ -VEX processor: Regular binary, Generic binary without optimizations and Generic binary with optimizations. These simulation have been performed with the XSTsim architectural simulator. The relative performance of generic binaries compared to regular binaries is displayed in Figure 5.6.




Figure 5.4: HP-LLVM relative performance




Figure 5.5: GCC-LLVM relative performance

Figure 5.6 shows that the optimization for generic binaries provides for more efficient generic binaries for most benchmarks. The `matrix` benchmark in particular shows excellent increase in performance. The `adpcm` benchmark however shows a significant decrease in performance. Closer inspection of the generated binaries shows that some benchmarks start spending a lot of time executing spill code to save registers to the stack. This is related to the optimization being too aggressive for `MachineBasicBlock` with a lot of virtual registers.

[13] stated that performance should increase if more registers are used for generic binaries. The amount of different registers that are used by each benchmark have been tracked and are displayed in 5.5. The table shows that the optimization is able to

Benchmark	W2	W4	W8
adpcm	2.206.300	2.115.820	2.118.040
bcnt	39.540	38.200	35.640
blit	1.164.280	1.204.140	1.204.120
compress	3.828.700	4.555.420	4.541.900
crc	1.323.160	1.290.540	1.290.560
engine	22.820.980	18.119.680	18.119.680
g3fax	55.537.540	56.645.740	56.680.260
matrix	717.100	615.060	615.040
pocsag	2.453.040	1.987.140	1.954.860
ucbqsort	9.568.680	9.457.160	9.454.600

Table 5.1: LLVM-based compiler performance in ns

Benchmark	W2	W4	W8	Comment
adpcm	763.120	613.520	ERR	Infinite loop
bcnt	16.860	11.100	11.080	
blit	341.620	301.360	301.360	
compress	2.893.080	2.421.420	2.400.140	
crc	486.980	347.040	347.040	
engine	15.785.220	13.720.900	13.719.860	
g3fax	24.821.120	22.678.220	ERR	Infinite loop
matrix	255.940	251.260	251.220	
pocsag	889.160	676.680	667.880	
ucbqsort	6.621.460	5.604.720	5.603.460	

Table 5.2: HP-based compiler performance in ns

increase the register usage for each benchmark. The `adpcm` benchmark already uses a large amount of registers and it is not possible to increase this by much. This probably

Benchmark	W4	Comment
adpcm	2.404.760	Wrong result
bcnt	ERR	Infinite loop
blit	1.284.340	Wrong result
compress	3.433.960	
crc	1.409.960	
engine	ERR	Error load
g3fax	62.666.300	
matrix	1.090.420	
pocsag	2.058.700	
ucbqsort	31.075.420	

Table 5.3: GCC-based compiler performance in ns




Figure 5.6: Generic-Regular performance

Benchmark	Regular	No optimization	With optimization
adpcm	176.327	187.135	194.492
blit	100.328	102.007	100.325
compress	385.797	392.115	391.920
crc	107.528	108.950	108.820
matrix	51.238	61.572	51.238
ucbqsort	165.578	172.152	163.860

Table 5.4: GCC-based compiler performance in ns

causes the excessive spill code that is generated for this benchmark.

5.4 Conclusion

In this section we have shown how the operation of the LLVM-based compiler has been verified and how well binaries execute that have been generated with the LLVM-based

Benchmark	Regular	No-opt	Opt
adpcm	49	49	56
blit	18	18	33
compress	27	27	60
crc	19	19	27
matrix	18	18	36
ucbqsort	24	24	31

Table 5.5: Register usage

compiler.

The benchmarks and verifications have shown that the LLVM-based compiler still contains bugs. Not all benchmarks are able to execute using the Modelsim simulator but all benchmarks are able to execute using the XSTsim simulator. This indicates that there are probably scheduling issues in the assembly that is generated.

We have shown that the LLVM-based compiler exceeds the performance of the GCC-based compiler but the compiler is still outperformed by the HP-based compiler. As expected the HP-based compiler generates binaries that perform very well. This is probably related to the trace based scheduling techniques that are employed to extract a high level of ILP. In addition the HP-based compiler also performs certain optimizations that are not available to the LLVM-based compiler at `-O0`.

Additionally the benchmarks have also shown that the LLVM-based compiler is the only compiler able to generate correct code for all selected benchmarks. Surprisingly, even the HP-based compiler generates incorrect binaries for certain benchmarks. The code quality of the GCC-based compiler is bad with four benchmarks failing to execute.

Further we have also shown that the generic binary optimization allows generic binaries to operate at speeds that are nearly equal to the regular binaries. The generic binary optimization does introduce spill code in benchmarks that already use a large number of physical registers.

Conclusion

6.1 Summary

In chapter 2 we discussed the background of the ρ -VEX processor and the basics of the LLVM compiler framework. Instructions that are supported by the ρ -VEX processor have been shown, the run-time architecture has been discussed and the register layout of the ρ -VEX processor has been discussed. The background information on the LLVM compiler should provide for a basic understanding on how the compiler operates and what parts need to be implemented to build a ρ -VEX backend.

Finally we have also discussed how the LLVM compiler will be verified. The verification step is extremely important to determine whether the binaries that are generated work correctly.

Chapter 3 discussed how the LLVM-based compiler has been implemented. A short description has been provided on how the Tablegen features are used for description of the backend. We have shown how the ρ -VEX processor has been described in Tablegen and how the different passes have been implemented. All the phases that are involved with code generation have been discussed.

Certain features that are required for the ρ -VEX processor are not available in the LLVM compiler. Features such as a machine description file are new to LLVM and we have shown what changes have been made to the LLVM compiler to support machine description files. In addition we have shown how the backend has been updated to provide support for the ρ -VEX generic binary format.

The chapter 4 discussed the optimizations that have been implemented to improve performance of the ρ -VEX binaries.

The `rvexMachineScheduler` pass is used to handle structural and data hazards. Temporary instruction packets are generated and are filled with instructions that have been selected using cost based scheduling algorithms to reduce register pressure. The pass enables ρ -VEX binaries to execute correctly and to perform better than binaries that have not been scheduled using the `rvexMachineScheduler`.

The branch analysis optimization is used to erase unnecessary `goto` statements from the code. In addition hooks have been provided that allow the LLVM `BranchFolding` pass to further optimize branches that are used in ρ -VEX binaries. The generic binary optimization allows binaries with generic binary support to perform on par with regular binaries. The performance of generic binaries will only degrade once the register pressure becomes too high and spill code needs to be inserted. The immediate value optimization allows more efficient use of available instructions of the ρ -VEX processor.

Finally in chapter 5 we have shown how the operation of the LLVM-based compiler has been verified and how well binaries execute that have been generated with the LLVM-based compiler.

The benchmarks and verifications have shown that the LLVM-based compiler still contains bugs. Not all benchmarks are able to execute using the Modelsim simulator but all benchmarks are able to execute using the XSTsim simulator. This indicates that there are probably scheduling issues in the assembly that is generated.

We have shown that the LLVM-based compiler exceeds the performance of the GCC-based compiler but the compiler is still outperformed by the HP-based compiler. As expected the HP-based compiler generates binaries that perform very well. This is probably related to the trace based scheduling techniques that are employed to extract a high level of ILP. In addition the HP-based compiler also performs certain optimizations that are not available to the LLVM-based compiler at -O0.

Additionally the benchmarks have also shown that the LLVM-based compiler is the only compiler able to generate correct code for all selected benchmarks. Surprisingly, even the HP-based compiler generates incorrect binaries for certain benchmarks. The code quality of the GCC-based compiler is bad with four benchmarks failing to execute.

Further we have also shown that the generic binary optimization allows generic binaries to operate at speeds that are nearly equal to the regular binaries. The generic binary optimization does introduce spill code in benchmarks that already use a large number of physical registers.

6.2 Main contributions

The following parts have been contributed to the ρ -VEX project:

- LLVM-based ρ -VEX compiler.
 - Floating point support.
 - 64-bit integer support
- Parameterization of LLVM-based compilers.
- Support for generic binaries.
- Scheduling optimizations that improve performance of regular binaries.
- Register allocation optimization that improves performance of generic binaries

6.3 Future work

Future work for the LLVM-based ρ -VEX could involve the following subjects.

- **Enhance parameterization:** At the moment only issue width, instruction stages and instruction delay can be customized through config files. In the future more configuration options can be added such as number of registers available or scheduling parameters.


- **LLVM JIT:** The LLVM compiler support Just-In-Time compilation through the LLVM interpreter. Implementing the interpreter for the rvex processor could produce interesting results where binary properties can be modified during runtime. For example the interpreter could look for code with higher degree of ILP. If suitable code is found the program will be executed on an 8-issue width rvex processor. If suitable code is not found the issue width could be reduced to 2- or 4-issue code and the idle functional units can be shut down to preserve energy.
- **Trace based scheduling:** Currently the LLVM compiler uses a basic block scheduler. Introducing a trace based scheduler could further improve performance of binaries for VLIW type processors.

Bibliography

- [1] D. A. P. John L. Hennessy, *Computer Organization and Design*. Morgan Kaufmann, 2009.
- [2] J. A. Fisher, *The VEX System*, pdf.
- [3] T. van As, “-vex: A reconfigurable and extensible vliw processor,” Master’s thesis, Technical university Delft, Mekelweg 4, 2628 CD, Delft, The Netherlands, 2008.
- [4] P. Faraboschi, G. Brown, J. Fisher, G. Desoll, and F. Homewood, “Lx: a technology platform for customizable vliw embedded processing,” in *Computer Architecture, 2000. Proceedings of the 27th International Symposium on*, 2000, pp. 203–213.
- [5] D. A. P. John L. Hennessy, *Computer Architecture, A Quantitative Approach, Fifth edition*. Morgan Kaufmann, 2012, ch. Chapter Three: Instruction-Level Parallelism and Its Exploitation, p. 244.
- [6] D. W. Wall, “Limits of instruction-level parallelism,” *WRL Research Report*, p. 73, November 1993.
- [7] J. A. Fisher, “Very long instruction word architectures and the eli-512,” *SIGARCH Comput. Archit. News*, vol. 11, no. 3, pp. 140–150, Jun. 1983. [Online]. Available: <http://dl.acm.org/citation.cfm?id=1067651.801649>
- [8] P. G. Lowney, P. G. Lowney, S. M. Freudenberger, T. J. Karzes, W. D. Lichtenstein, R. P. Nix, J. S. O’donnell, and J. C. Ruttenberg, “The multiflow trace scheduling compiler,” *JOURNAL OF SUPERCOMPUTING*, vol. 7, pp. 51–142, 1993.
- [9] C. Y. Joseph A. Fisher, Paolo Faraboschi, *Embedded computing: A VLIW Approach to Architecture, Compilers and Tools*. Elsevier, 2005.
- [10] R. Seedorf, “Fingerprint verification on the vex processor,” Master’s thesis, Technical university Delft, 2011.
- [11] Clang, “Clang - features and goals.”
- [12] C. Lattner, “The llvm compiler framework and infrastructure (part 1),” Presentation.
- [13] S. W. Anthony Brandon, “Support for dynamic issue width in vliw processors using generic binaries,” *EDAA*, 2013.
- [14] J. A. B. M. Jeff Scott, Lea Hwang Lee, “Designing the low-power m*core architecture,” in *Proc. IEEE Power Driven Microarchitecture Workshop*, 1998.

List of definitions

.. ...



asd

A.1 LLVM toolchain

asd

A.2 XSTsim

asd

asd

B.1 Building LLVM from source

asd